

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: df = pd.read_csv("Social_Network_Ads.csv")
```

```
In [3]: #check the head
df.head()
```

Out[3]:

| | User ID | Gender | Age | EstimatedSalary | Purchased |
|---|----------|--------|-----|-----------------|-----------|
| 0 | 15624510 | Male | 19 | 19000 | 0 |
| 1 | 15810944 | Male | 35 | 20000 | 0 |
| 2 | 15668575 | Female | 26 | 43000 | 0 |
| 3 | 15603246 | Female | 27 | 57000 | 0 |
| 4 | 15804002 | Male | 19 | 76000 | 0 |

```
In [4]: #Check tail
df.tail()
```

Out[4]:

| | User ID | Gender | Age | EstimatedSalary | Purchased |
|-----|----------|--------|-----|-----------------|-----------|
| 395 | 15691863 | Female | 46 | 41000 | 1 |
| 396 | 15706071 | Male | 51 | 23000 | 1 |
| 397 | 15654296 | Female | 50 | 20000 | 1 |
| 398 | 15755018 | Male | 36 | 33000 | 0 |
| 399 | 15594041 | Female | 49 | 36000 | 1 |

```
In [5]: df.dtypes
```

```
Out[5]: User ID          int64
Gender          object
Age            int64
EstimatedSalary int64
Purchased       int64
dtype: object
```

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User ID               400 non-null   int64
1   Gender                400 non-null   object
2   Age                   400 non-null   int64
3   EstimatedSalary       400 non-null   int64
4   Purchased             400 non-null   int64
dtypes: int64(4), object(1)
memory usage: 15.8+ KB
```

In [7]: `df.describe()`

Out[7]:

| | User ID | Age | EstimatedSalary | Purchased |
|--------------|--------------|------------|-----------------|------------|
| count | 4.000000e+02 | 400.000000 | 400.000000 | 400.000000 |
| mean | 1.569154e+07 | 37.655000 | 69742.500000 | 0.357500 |
| std | 7.165832e+04 | 10.482877 | 34096.960282 | 0.479864 |
| min | 1.556669e+07 | 18.000000 | 15000.000000 | 0.000000 |
| 25% | 1.562676e+07 | 29.750000 | 43000.000000 | 0.000000 |
| 50% | 1.569434e+07 | 37.000000 | 70000.000000 | 0.000000 |
| 75% | 1.575036e+07 | 46.000000 | 88000.000000 | 1.000000 |
| max | 1.581524e+07 | 60.000000 | 150000.000000 | 1.000000 |

In [8]: *#checking the duplicacy*
`df.duplicated().sum()`

Out[8]: 0

In [9]: `df.isnull().sum()`

Out[9]: User ID 0
Gender 0
Age 0
EstimatedSalary 0
Purchased 0
dtype: int64

```
In [10]: df.corr()
```

```
Out[10]:
```


| | User ID | Age | EstimatedSalary | Purchased |
|-----------------|-----------|-----------|-----------------|-----------|
| User ID | 1.000000 | -0.000721 | 0.071097 | 0.007120 |
| Age | -0.000721 | 1.000000 | 0.155238 | 0.622454 |
| EstimatedSalary | 0.071097 | 0.155238 | 1.000000 | 0.362083 |
| Purchased | 0.007120 | 0.622454 | 0.362083 | 1.000000 |

```
In [11]: df["Gender"].value_counts(normalize = True)*100
```

```
Out[11]: Female    51.0
Male        49.0
Name: Gender, dtype: float64
```


```
In [12]: sns.countplot(x = "Gender",data = df)
```

```
Out[12]: <AxesSubplot:xlabel='Gender', ylabel='count'>
```



```
In [13]: df["Gender"].value_counts().plot(kind = "pie", autopct = "%.2f")
```

```
Out[13]: <AxesSubplot:ylabel='Gender'>
```



```
In [14]: def age_group(pi):
    if pi > 20 and pi <= 30:
        Age_group = "21-30"
    elif pi > 30 and pi <= 40:
        Age_group = "31-40"
    elif pi > 40 and pi <= 50:
        Age_group = "41-50"
    else:
        Age_group = "Above 50"
    return (Age_group)
```

```
In [15]: df["Age_group"] = df["Age"].apply(age_group)
```

```
In [16]: df
```

```
Out[16]:
```

| | User ID | Gender | Age | EstimatedSalary | Purchased | Age_group |
|-----|----------|--------|-----|-----------------|-----------|-----------|
| 0 | 15624510 | Male | 19 | 19000 | 0 | Above 50 |
| 1 | 15810944 | Male | 35 | 20000 | 0 | 31-40 |
| 2 | 15668575 | Female | 26 | 43000 | 0 | 21-30 |
| 3 | 15603246 | Female | 27 | 57000 | 0 | 21-30 |
| 4 | 15804002 | Male | 19 | 76000 | 0 | Above 50 |
| ... | ... | ... | ... | ... | ... | ... |
| 395 | 15691863 | Female | 46 | 41000 | 1 | 41-50 |
| 396 | 15706071 | Male | 51 | 23000 | 1 | Above 50 |
| 397 | 15654296 | Female | 50 | 20000 | 1 | 41-50 |
| 398 | 15755018 | Male | 36 | 33000 | 0 | 31-40 |
| 399 | 15594041 | Female | 49 | 36000 | 1 | 41-50 |


400 rows × 6 columns

```
In [17]: df.columns
```

```
Out[17]: Index(['User ID', 'Gender', 'Age', 'EstimatedSalary', 'Purchased',
               'Age_group'],
              dtype='object')
```

```
In [18]: sns.countplot(data = df, x = "Age_group")
```

```
Out[18]: <AxesSubplot:xlabel='Age_group', ylabel='count'>
```



```
In [19]: pd.crosstab(df["Age_group"],df["Purchased"],margins=True)
```

```
Out[19]:
```


| Purchased | 0 | 1 | All |
|-----------|-----|-----|-----|
| Age_group | | | |
| 21-30 | 86 | 6 | 92 |
| 31-40 | 109 | 33 | 142 |
| 41-50 | 39 | 59 | 98 |
| Above 50 | 23 | 45 | 68 |
| All | 257 | 143 | 400 |

```
In [20]: sns.distplot(df["Age"])
```

C:\Users\HP\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


```
warnings.warn(msg, FutureWarning)
```

```
Out[20]: <AxesSubplot:xlabel='Age', ylabel='Density'>
```

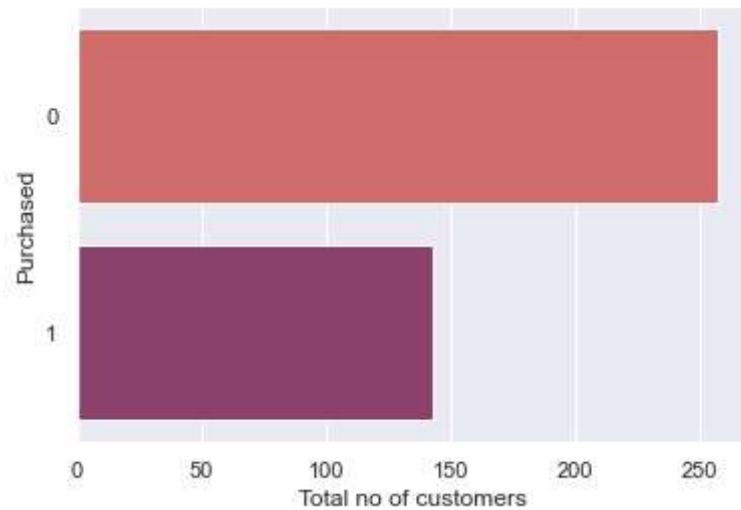


```
In [21]: sns.catplot(data=df, y="Purchased", x="Age_group", kind="bar")
```


```
Out[21]: <seaborn.axisgrid.FacetGrid at 0x275ca99cb50>
```



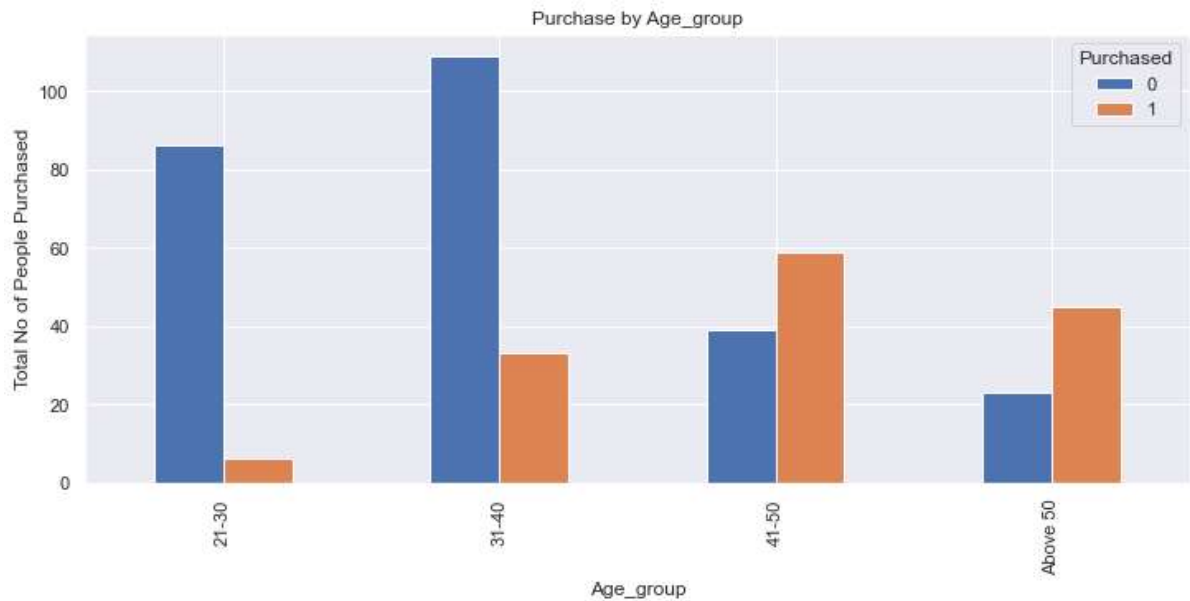
```
In [22]: sns.set_theme(style="darkgrid")
sns.countplot(y="Purchased",data=df, palette = "flare")
plt.ylabel("Purchased")
plt.xlabel("Total no of customers")
plt.show()
```



```
In [23]: sns.set_theme(style = "darkgrid")
sns.countplot(data=df,x="Gender", palette = "rocket")
plt.xlabel("Gender")
plt.ylabel("Total customers")
plt.show()
```




```
In [24]: pd.crosstab(df["Age_group"],df["Purchased"]).plot(kind="bar",figsize=(12,5))
plt.title("Purchase by Age_group")
plt.xlabel("Age_group")
plt.ylabel("Total No of People Purchased")
plt.show()
```



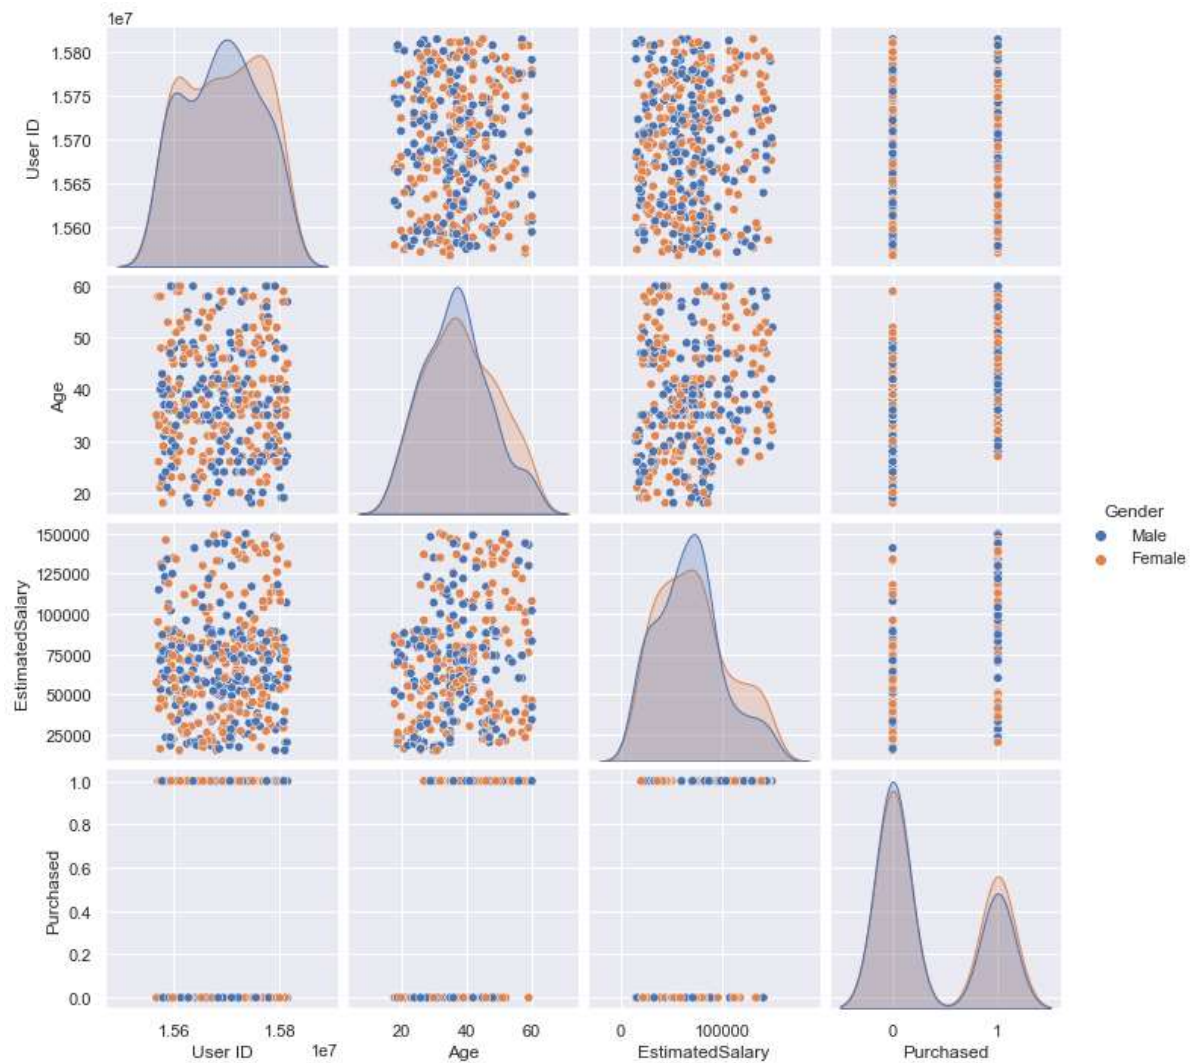
```
In [25]: sns.jointplot(data = df,x="EstimatedSalary",y="Purchased",hue = "Gender",kind
```

```
Out[25]: <seaborn.axisgrid.JointGrid at 0x275caf246d0>
```




```
In [26]: sns.pairplot(data = df,hue="Gender")
```

```
Out[26]: <seaborn.axisgrid.PairGrid at 0x275cb071c40>
```



In [27]: df

Out[27]:

| | User ID | Gender | Age | EstimatedSalary | Purchased | Age_group |
|-----|----------|--------|-----|-----------------|-----------|-----------|
| 0 | 15624510 | Male | 19 | 19000 | 0 | Above 50 |
| 1 | 15810944 | Male | 35 | 20000 | 0 | 31-40 |
| 2 | 15668575 | Female | 26 | 43000 | 0 | 21-30 |
| 3 | 15603246 | Female | 27 | 57000 | 0 | 21-30 |
| 4 | 15804002 | Male | 19 | 76000 | 0 | Above 50 |
| ... | ... | ... | ... | ... | ... | ... |
| 395 | 15691863 | Female | 46 | 41000 | 1 | 41-50 |
| 396 | 15706071 | Male | 51 | 23000 | 1 | Above 50 |
| 397 | 15654296 | Female | 50 | 20000 | 1 | 41-50 |
| 398 | 15755018 | Male | 36 | 33000 | 0 | 31-40 |
| 399 | 15594041 | Female | 49 | 36000 | 1 | 41-50 |

400 rows × 6 columns

In [28]: x = df.iloc[:,2:4].values

In [29]: y = df.iloc[:,5].values

```
In [30]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.3 , random_

from sklearn.neighbors import KNeighborsClassifier
KNClassifier = KNeighborsClassifier(n_neighbors = 17)
KNClassifier.fit(x_train,y_train)
from sklearn.metrics import classification_report

y_pred = KNClassifier.predict(x_test)

print(classification_report(y_test,y_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.85 | 0.97 | 0.91 | 79 |
| 1 | 0.93 | 0.66 | 0.77 | 41 |
| accuracy | | | 0.87 | 120 |
| macro avg | 0.89 | 0.82 | 0.84 | 120 |
| weighted avg | 0.88 | 0.87 | 0.86 | 120 |

```
In [31]: from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
KNacc = accuracy_score(y_pred,y_test)
print(confusion_matrix(y_test,y_pred))
print("K neighbors Accuracy score is: {:.2f}%".format(KNacc*100))
```

```
[[77  2]
 [14 27]]
K neighbors Accuracy score is: 86.67%
```

```
In [ ]:
```